

# Classifying Indian denomiation

## Importing libraries

In [1]:

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 import os
4 import PIL
5 import tensorflow as tf
6
7 from tensorflow import keras
8 from tensorflow.keras import layers
9 from tensorflow.keras.models import Sequential
```

## Setting the path to the dataset

In [2]:

```
1 import pathlib
2 data_dir = '/home/glitch101/Downloads/CURRENCY_DATASET'
3 data_dir = pathlib.Path(data_dir)
```

## Checking the number of images in the dataset

In [3]:

```
1 image_count = len(list(data_dir.glob('*/*.png')))
2 print(image_count)
```

400

## Samples of each class

In [4]:

```
1 ten = list(data_dir.glob('10/*'))  
2 PIL.Image.open(str(ten[0]))
```



In [5]:

1 PIL.Image.open(str(ten[1]))



In [6]:

```
1 twenty = list(data_dir.glob('20/*'))  
2 PIL.Image.open(str(twenty[0]))
```



In [7]:

1 PIL.Image.open(str(twenty[1]))



In [8]:

```
1 hundred = list(data_dir.glob('100/*'))  
2 PIL.Image.open(str(hundred[0]))
```



In [9]:

```
1 PIL.Image.open(str(hundred[1]))
```



In [10]:

```
1 two_hundred = list(data_dir.glob('200/*'))  
2 PIL.Image.open(str(two_hundred[0]))
```



In [11]:

```
1 PIL.Image.open(str(two_hundred[1]))
```



## Defining some parameters for the loader

In [12]:

```
1 batch_size = 32
2 img_height = 180
3 img_width = 180
```

I have used a 80-20 ratio for training and testing images

In [13]:

```
1 train_ds = tf.keras.preprocessing.image_dataset_from_directory(  
2     data_dir,  
3     validation_split=0.2,  
4     subset="training",  
5     seed=123,  
6     image_size=(img_height, img_width),  
7     batch_size=batch_size)
```

Found 556 files belonging to 4 classes.  
Using 445 files for training.

In [14]:

```
1 val_ds = tf.keras.preprocessing.image_dataset_from_directory(  
2     data_dir,  
3     validation_split=0.2,  
4     subset="validation",  
5     seed=123,  
6     image_size=(img_height, img_width),  
7     batch_size=batch_size)
```

Found 556 files belonging to 4 classes.  
Using 111 files for validation.

## We have four classes 10, 100, 20, 200

In [15]:

```
1 class_names = train_ds.class_names  
2 print(class_names)
```

['10', '100', '20', '200']

## Visualizing the data

In [16]:

```
1 import matplotlib.pyplot as plt
2
3 plt.figure(figsize=(10, 10))
4 for images, labels in train_ds.take(1):
5     for i in range(9):
6         ax = plt.subplot(3, 3, i + 1)
7         plt.imshow(images[i].numpy().astype("uint8"))
8         plt.title(class_names[labels[i]])
9         plt.axis("off")
```



## Creating a batch

```
In [17]:  
1 for image_batch, labels_batch in train_ds:  
2     print(image_batch.shape)  
3     print(labels_batch.shape)  
4     break  
  
(32, 180, 180, 3)  
(32,)
```

## Configuration

```
In [18]:  
1 AUTOTUNE = tf.data.experimental.AUTOTUNE  
2  
3 train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)  
4 val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

## Rescaling the rgb channels from [0, 255] to [0, 1]

```
In [19]:  
1 normalization_layer = layers.experimental.preprocessing.Rescaling(1./255)
```

```
In [20]:  
1 normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))  
2 image_batch, labels_batch = next(iter(normalized_ds))  
3 first_image = image_batch[0]  
4 # Notice the pixels values are now in `[0,1]`.  
5 print(np.min(first_image), np.max(first_image))
```

0.0028860057 1.0

## Creating the model

In [21]:

```
1 num_classes = 4
2
3 model = Sequential([
4     layers.experimental.preprocessing.Rescaling(1./255, input_shape=(img_heig
5     layers.Conv2D(16, 3, padding='same', activation='relu'),
6     layers.MaxPooling2D(),
7     layers.Conv2D(32, 3, padding='same', activation='relu'),
8     layers.MaxPooling2D(),
9     layers.Conv2D(64, 3, padding='same', activation='relu'),
10    layers.MaxPooling2D(),
11    layers.Flatten(),
12    layers.Dense(128, activation='relu'),
13    layers.Dense(num_classes)
14])
```

---

## Compiling the model

---

In [22]:

```
1 model.compile(optimizer='adam',
2                 loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
3                 metrics=['accuracy'])
```

---

In [23]:

1 model.summary()

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
rescaling_1 (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d (MaxPooling2D)	(None, 90, 90, 16)	0
conv2d_1 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 45, 45, 32)	0
conv2d_2 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 22, 22, 64)	0
flatten (Flatten)	(None, 30976)	0
dense (Dense)	(None, 128)	3965056
dense_1 (Dense)	(None, 4)	516
<hr/>		
Total params: 3,989,156		
Trainable params: 3,989,156		
Non-trainable params: 0		

## Training

In [24]:

```
1 epochs=10
2 history = model.fit(
3     train_ds,
4     validation_data=val_ds,
5     epochs=epochs
6 )
```

```
Epoch 1/10
14/14 [=====] - 7s 493ms/step - loss: 1.3468 - accuracy: 0.3933 - val_loss: 0.5045
Epoch 2/10
14/14 [=====] - 5s 341ms/step - loss: 0.8527 - accuracy: 0.6584 - val_loss: 0.6757
Epoch 3/10
14/14 [=====] - 5s 352ms/step - loss: 0.5009 - accuracy: 0.8022 - val_loss: 0.7027
Epoch 4/10
14/14 [=====] - 5s 383ms/step - loss: 0.3507 - accuracy: 0.8697 - val_loss: 0.8108
Epoch 5/10
14/14 [=====] - 5s 325ms/step - loss: 0.1965 - accuracy: 0.9258 - val_loss: 0.7838
Epoch 6/10
14/14 [=====] - 7s 489ms/step - loss: 0.1350 - accuracy: 0.9573 - val_loss: 0.8559
Epoch 7/10
14/14 [=====] - 6s 456ms/step - loss: 0.0913 - accuracy: 0.9685 - val_loss: 0.7658
Epoch 8/10
14/14 [=====] - 7s 481ms/step - loss: 0.0836 - accuracy: 0.9775 - val_loss: 0.8108
Epoch 9/10
14/14 [=====] - 7s 474ms/step - loss: 0.0355 - accuracy: 0.9955 - val_loss: 0.8018
Epoch 10/10
14/14 [=====] - 6s 440ms/step - loss: 0.0267 - accuracy: 0.9955 - val_loss: 0.8288
```

---

## Training results

---

In [25]:

```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3
4 loss = history.history['loss']
5 val_loss = history.history['val_loss']
6
7 epochs_range = range(epochs)
8
9 plt.figure(figsize=(8, 8))
10 plt.subplot(1, 2, 1)
11 plt.plot(epochs_range, acc, label='Training Accuracy')
12 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
13 plt.legend(loc='lower right')
14 plt.title('Training and Validation Accuracy')
15
16 plt.subplot(1, 2, 2)
17 plt.plot(epochs_range, loss, label='Training Loss')
18 plt.plot(epochs_range, val_loss, label='Validation Loss')
19 plt.legend(loc='upper right')
20 plt.title('Training and Validation Loss')
21 plt.show()
```



## Augmenting the dataset to improvise for small training samples

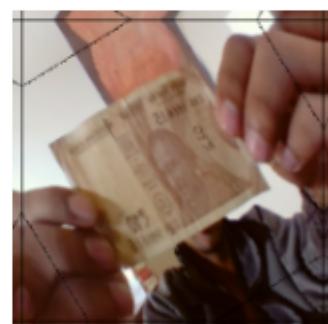
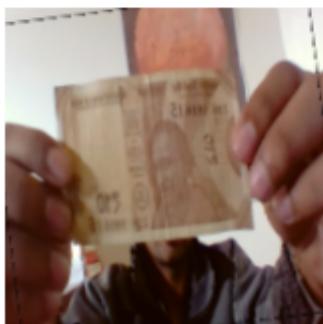
In [26]:

```
1 data_augmentation = keras.Sequential(  
2     [  
3         layers.experimental.preprocessing.RandomFlip("horizontal",  
4             input_shape=(img_height,  
5                 img_width,  
6                     3)),  
7         layers.experimental.preprocessing.RandomRotation(0.1),  
8         layers.experimental.preprocessing.RandomZoom(0.1),  
9     ]  
10 )
```

## Few augmentaiton examples

In [27]:

```
1 plt.figure(figsize=(10, 10))
2 for images, _ in train_ds.take(1):
3     for i in range(9):
4         augmented_images = data_augmentation(images)
5         ax = plt.subplot(3, 3, i + 1)
6         plt.imshow(augmented_images[0].numpy().astype("uint8"))
7         plt.axis("off")
```



Using Dropout to reduce overfitting of the model

In [28]:

```
1 model = Sequential([
2     data_augmentation,
3     layers.experimental.preprocessing.Rescaling(1./255),
4     layers.Conv2D(16, 3, padding='same', activation='relu'),
5     layers.MaxPooling2D(),
6     layers.Conv2D(32, 3, padding='same', activation='relu'),
7     layers.MaxPooling2D(),
8     layers.Conv2D(64, 3, padding='same', activation='relu'),
9     layers.MaxPooling2D(),
10    layers.Dropout(0.2),
11    layers.Flatten(),
12    layers.Dense(128, activation='relu'),
13    layers.Dense(num_classes)
14])
```

---

## Compiling the model

---

In [29]:

```
1 model.compile(optimizer='adam',
2                 loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
3                 metrics=['accuracy'])
```

In [30]:

1 model.summary()

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
<hr/>		
sequential_1 (Sequential)	(None, 180, 180, 3)	0
rescaling_2 (Rescaling)	(None, 180, 180, 3)	0
conv2d_3 (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d_3 (MaxPooling2D)	(None, 90, 90, 16)	0
conv2d_4 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_4 (MaxPooling2D)	(None, 45, 45, 32)	0
conv2d_5 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_5 (MaxPooling2D)	(None, 22, 22, 64)	0
dropout (Dropout)	(None, 22, 22, 64)	0
flatten_1 (Flatten)	(None, 30976)	0
dense_2 (Dense)	(None, 128)	3965056
dense_3 (Dense)	(None, 4)	516
<hr/>		
Total params: 3,989,156		
Trainable params: 3,989,156		
Non-trainable params: 0		

## Training the model

In [31]:

```
1 epochs = 200
2 history = model.fit(
3     train_ds,
4     validation_data=val_ds,
5     epochs=epochs
6 )

14/14 [=====] - 5s 367ms/step - loss: 0.0012 - accuracy: 1.0000 - val_loss: 0.9279
Epoch 167/200
14/14 [=====] - 5s 365ms/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 0.9369
Epoch 168/200
14/14 [=====] - 5s 377ms/step - loss: 0.0011 - accuracy: 1.0000 - val_loss: 0.9279
Epoch 169/200
14/14 [=====] - 5s 368ms/step - loss: 5.6975e-04 - accuracy: 1.0000 - val_loss: 0.9369
Epoch 170/200
14/14 [=====] - 5s 368ms/step - loss: 0.0011 - accuracy: 1.0000 - val_loss: 0.9279
Epoch 171/200
```

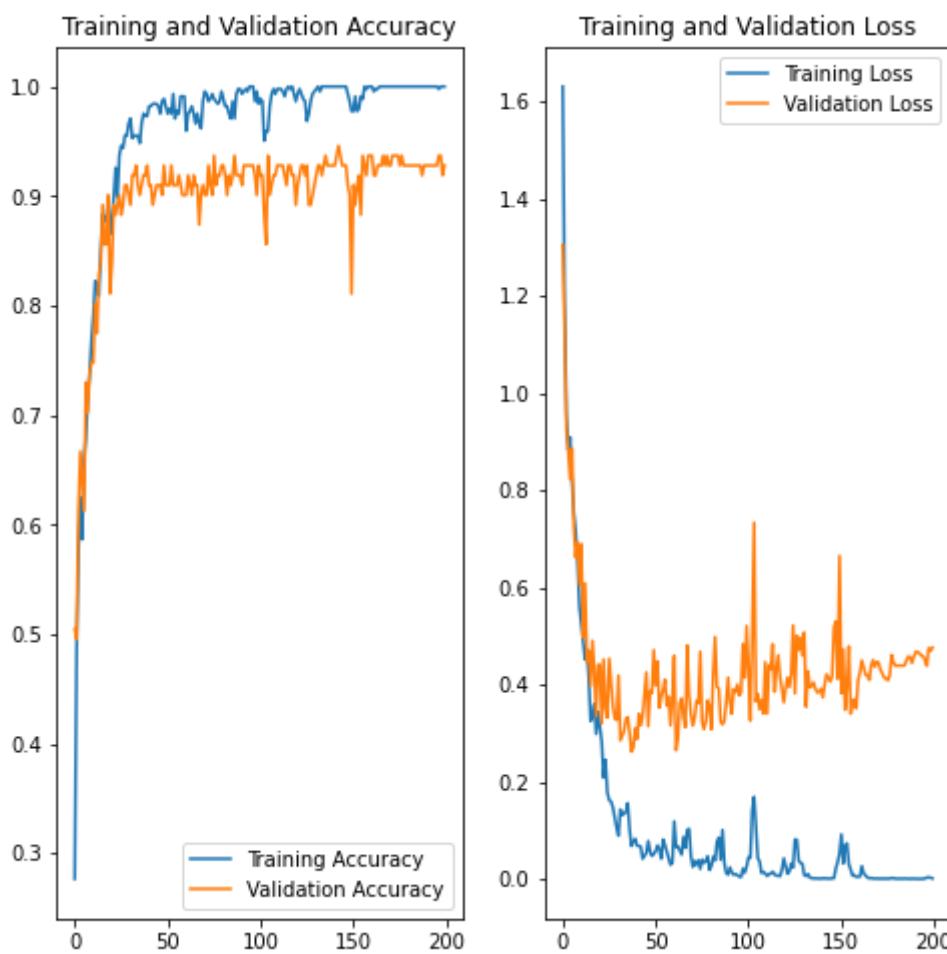
---

## Training results

---

In [35]:

```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3
4 loss = history.history['loss']
5 val_loss = history.history['val_loss']
6
7 epochs_range = range(epochs)
8
9 plt.figure(figsize=(8, 8))
10 plt.subplot(1, 2, 1)
11 plt.plot(epochs_range, acc, label='Training Accuracy')
12 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
13 plt.legend(loc='lower right')
14 plt.title('Training and Validation Accuracy')
15
16 plt.subplot(1, 2, 2)
17 plt.plot(epochs_range, loss, label='Training Loss')
18 plt.plot(epochs_range, val_loss, label='Validation Loss')
19 plt.legend(loc='upper right')
20 plt.title('Training and Validation Loss')
21 plt.show()
```



## Testing on new data

In [36]:

```
1 test_image_path = '/home/glitch101/Downloads/20_rupee.jpg'
2
3 img = keras.preprocessing.image.load_img(
4     test_image_path, target_size=(img_height, img_width)
5 )
6 img_array = keras.preprocessing.image.img_to_array(img)
7 img_array = tf.expand_dims(img_array, 0) # Create a batch
8
9 predictions = model.predict(img_array)
10 score = tf.nn.softmax(predictions[0])
11
12 print(
13     "This image most likely belongs to {} with a {:.2f} percent confidence."
14     .format(class_names[np.argmax(score)], 100 * np.max(score))
15 )
16
17 print(score)
```

This image most likely belongs to 20 with a 100.00 percent confidence.  
tf.Tensor([5.7689140e-06 1.9313859e-10 9.9999404e-01 2.7660934e-07], shape=(4,), dtype=float32)

## Saving the model for use with our python script

In [37]:

```
1 model.save('/home/glitch101/Downloads/CURRENCY_MODEL/')
```

INFO:tensorflow:Assets written to: /home/glitch101/Downloads/CURRENCY\_MODEL/assets